

Global Descriptor Table

From OSDev Wiki

The **Global Descriptor Table (GDT)** is a binary data structure specific to the IA-32 and x86-64 architectures. It contains entries telling the CPU about memory segments. A similar Interrupt Descriptor Table exists containing task and interrupt descriptors.

It is recommended to read the GDT Tutorial.

Contents

- 1 GDTR
- 2 Table
- 3 Segment Descriptor
- 4 System Segment Descriptor
- 5 Long Mode System Segment Descriptor
- 6 See Also
 - 6.1 Articles
 - 6.2 External references

GDTR

The **GDT** is pointed to by the value in the **GDTR** register. This is loaded using the **LGDT** assembly instruction, whose argument is a pointer to a **GDT Descriptor** structure:

GDT Descriptor (GDTR)

79 (64-bit mode) 48 (32-bit mode)	16	15	0
Offset 63 (64-bit mode) 31 (32-bit mode)	Size 0	15	0

- **Size:** The size of the table in bytes subtracted by 1. This subtraction occurs because the maximum value of **Size** is 65535, while the **GDT** can be up to 65536 bytes in length (8192 entries). Further, no **GDT** can have a size of 0 bytes.
- **Offset:** The linear address of the **GDT** (not the physical address, paging applies).

Note that the amount of data loaded by **LGDT** differs in 32-bit and 64-bit modes, the offset is 4 bytes long in 32-bit mode and 8 bytes long in 64-bit mode.

For more information, see **Section 2.4.1: Global Descriptor Table Register (GDTR)** and **Figure 2-6: Memory Management Registers** of the Intel Software

Table

The entries in the **GDT** are 8 bytes long and form a table like this:

Global Descriptor Table	
Address	Content
GDTR Offset + 0	Null
GDTR Offset + 8	Entry 1
GDTR Offset + 16	Entry 2
GDTR Offset + 24	Entry 3
...	...

The first entry in the **GDT** (Entry 0) should always be null and subsequent entries should be used instead.

Entries in the table are accessed by **Segment Selectors**, which are loaded into **Segmentation** registers either by assembly instructions or by hardware functions such as **Interrupts**.

Segment Descriptor

Each entry in the table has a complex structure:

Segment Descriptor										
63	56	55	52	51	48	47	40	39	32	
Base		Flags		Limit		Access Byte		Base		
31	24	3	0	19	16	7	0	23	16	
31					16		15			0
Base						Limit				
15						0 15				0

- **Base:** A 32-bit value containing the linear address where the segment begins.
- **Limit:** A 20-bit value, tells the maximum addressable unit, either in 1 byte units, or in 4KiB pages. Hence, if you choose page granularity and set the **Limit** value to 0xFFFFF the segment will span the full 4 GiB address space in 32-bit mode.

In 64-bit mode, the **Base** and **Limit** values are ignored, each descriptor covers the entire linear address space regardless of what they are set to.

For more information, see **Section 3.4.5: Segment Descriptors** and **Figure 3-8: Segment Descriptor** of the Intel Software Developer Manual, Volume 3-A.

Access Byte							
7	6	5	4	3	2	1	0
P	DPL		S	E	DC	RW	A

- **P:** Present bit. Allows an entry to refer to a valid segment. Must be set (**1**) for

any valid segment.

- **DPL:** Descriptor privilege level field. Contains the CPU Privilege level of the segment. **0** = highest privilege (kernel), **3** = lowest privilege (user applications).
- **S:** Descriptor type bit. If clear (**0**) the descriptor defines a system segment (eg. a Task State Segment). If set (**1**) it defines a code or data segment.
- **E:** Executable bit. If clear (**0**) the descriptor defines a data segment. If set (**1**) it defines a code segment which can be executed from.
- **DC:** Direction bit/Conforming bit.
 - For data selectors: Direction bit. If clear (**0**) the segment grows up. If set (**1**) the segment grows down, ie. the **Offset** has to be greater than the **Limit**.
 - For code selectors: Conforming bit.
 - If clear (**0**) code in this segment can only be executed from the ring set in **DPL**.
 - If set (**1**) code in this segment can be executed from an equal or lower privilege level. For example, code in ring 3 can far-jump to *conforming* code in a ring 2 segment. The **DPL** field represent the highest privilege level that is allowed to execute the segment. For example, code in ring 0 cannot far-jump to a conforming code segment where **DPL** is 2, while code in ring 2 and 3 can. Note that the privilege level remains the same, ie. a far-jump from ring 3 to a segment with a **DPL** of 2 remains in ring 3 after the jump.
- **RW:** Readable bit/Writable bit.
 - For code segments: Readable bit. If clear (**0**), read access for this segment is not allowed. If set (**1**) read access is allowed. Write access is never allowed for code segments.
 - For data segments: Writeable bit. If clear (**0**), write access for this segment is not allowed. If set (**1**) write access is allowed. Read access is always allowed for data segments.
- **A:** Accessed bit. Best left clear (**0**), the CPU will set it when the segment is accessed.

Flags

3	2	1	0
G	DB	L	Reserved

- **G:** Granularity flag, indicates the size the **Limit** value is scaled by. If clear (**0**), the **Limit** is in 1 Byte blocks (byte granularity). If set (**1**), the **Limit** is in 4 KiB blocks (page granularity).
- **DB:** Size flag. If clear (**0**), the descriptor defines a 16-bit protected mode segment. If set (**1**) it defines a 32-bit protected mode segment. A GDT can have both 16-bit and 32-bit selectors at once.
- **L:** Long-mode code flag. If set (**1**), the descriptor defines a 64-bit code segment. When set, **DB** should always be clear. For any other type of segment (other code types or any data segment), it should be clear (**0**).

System Segment Descriptor

For system segments, such as those defining a **Task State Segment** or **Local Descriptor Table**, the format of the **Access Byte** differs slightly, in order to define different types of system segments rather than code and data segments.

For more information, see **Section 3.5: System Descriptor Types** and **Figure 3-2: System-Segment and Gate-Descriptor Types** of the Intel Software Developer Manual, Volume 3-A.

Access Byte							
7	6	5	4	3	2	1	0
P		DPL		S	Type		

- **Type:** Type of system segment.

Types available in 32-bit protected mode:

- **0x1:** 16-bit TSS (Available)
- **0x2:** LDT
- **0x3:** 16-bit TSS (Busy)
- **0x9:** 32-bit TSS (Available)
- **0xB:** 32-bit TSS (Busy)

Types available in Long Mode:

- **0x2:** LDT
- **0x9:** 64-bit TSS (Available)
- **0xB:** 64-bit TSS (Busy)

Long Mode System Segment Descriptor

For a **Task State Segment** or **Local Descriptor Table** in **Long Mode**, the format of a **Segment Descriptor** differs to ensure that the **Base** value can contain a 64-bit **Linear Address**. It takes up the space in the table of two usual entries, in a little endian format, such that the lower half of this entry precedes the higher half in the table.

For more information, see **Section 7.2.3: TSS Descriptor in 64-bit Mode** and **Figure 7-4: Format of TSS and LDT Descriptors in 64-bit Mode** of the Intel Software Developer Manual, Volume 3-A.

64-bit System Segment Descriptor									
127									96
Reserved									
95									64
Base 63									32
63	56	55	52	51	48	47	40	39	32
Base 3124		Flags 30		Limit 1916		Access Byte 70		Base 2316	
31					16		150		
Base 150						Limit 150			

See Also

Articles

- GDT Tutorial
- Getting to Ring 3
- Segmentation
- <http://www.osdever.net/tutorials/view/the-world-of-protected-mode> - how to set up GDT in assembler

External references

- Protected Mode tutorial (<http://files.osdev.org/mirrors/geezer/os/pm.htm>)
- Intel® 64 and IA-32 Architectures Software Developer's Manual, Volume 3A:. System Programming Guide, Part 1 (order number 253668) (<http://www.intel.com/design/processor/manuals/253668.pdf>) chapter 2.4

Retrieved from "https://wiki.osdev.org/index.php?title=Global_Descriptor_Table&oldid=27067"

Category: X86 CPU

-
- This page was last modified on 4 May 2022, at 05:38.
 - This page has been accessed 329,961 times.